

The Problem

Detecting coordinated inauthentic account networks is expensive with content-based approaches: they require direct platform API access, labeled training sets, and break immediately when the coordinated narrative shifts. Platforms typically catch campaigns weeks or months after activity peaks.

The Approach

GhostSig fingerprints public behavioral metadata — posting timing cadence, cross-platform activity rhythm, linguistic entropy, and device-echo signals — using an end-to-end ML pipeline: neural encoders (temporal, entropy, fusion) feed HDBSCAN clustering, which is scored by an XGBoost adversarial discriminator trained on synthetic bot/organic data. No content is read. No PII is collected.

Verified Results

Metric	Value	Source
Automated tests passing	44 / 44	tests/ (pytest)
Largest detected campaign (accounts)	57 accounts · confidence 0.9931	docs/evidence/*.json
Second-largest campaign	33 accounts · confidence 0.9931	docs/evidence/*.json
Campaigns detected (all HIGH tier)	12 campaigns	docs/evidence/*.json
REST API endpoints	10	api/main.py
Evidence output formats	JSON + PDF per campaign	attribution/evidence_cards.py

This is an analyst-support signal, not an automated enforcement decision. Human review required.

What I'm Looking For

Pilot feedback from a Trust & Safety or OSINT team, a labeled CIB dataset to validate against real-world campaigns, or a 15-minute call.